





**Say it once. Every screen knows.**

**Voice control for restaurant kitchens, built on the AssemblyAI Voice Agent API.**

A cook says "86 salmon" → every screen updates in seconds.

<https://86-it.rd5510.workers.dev/station>

## The problem

- When a kitchen runs out of a dish, **the news travels by shouting**.
- Servers keep selling it — so do DoorDash and Uber Eats.
- Result: refunds, comped meals, angry guests, 1-star reviews.
- Cooks **can't tap screens** with gloved or greasy hands.
- Kitchens are **too loud** for ordinary voice assistants.


 Add 1–2 sourced numbers here (refund rates, razor-thin restaurant margins) before submitting.

# The demo (live)

| The cook says                          | What happens                                                           |
|----------------------------------------|------------------------------------------------------------------------|
| "86 salmon"                            | Salmon greys out on KDS, servers' screen and online menu — same moment |
| "Fire 12" / "Bump 4"                   | Ticket moves on the kitchen display                                    |
| "How long on 7?"                       | "Seven: 14 minutes."                                                   |
| "12 has a nut allergy"                 | Red banner on ticket 12                                                |
| "Void 7?" → "yes"                      | Confirmed before anything happens                                      |
| <i>(teammate chats about football)</i> | <b>Nothing</b> — the agent stays silent                                |
| <i>(ticket passes 10 min)</i>          | Agent speaks up: "Table 4, 12 minutes."                                |

Live: <https://86-it.rd5510.workers.dev/station>

## How it works



**AssemblyAI features used:** Voice Agent API · inline session config · 10 client tools with JSON-Schema enums · keyterms (kitchen slang) · voice focus `near-field` · `transcription_mode: min_latency` · semantic barge-in · temporary tokens · `reply.create` alerts · `session.resume` reconnect

# Results

43-clip eval: 31 commands + 12 chatter lines, each in a fresh session.

| Condition         | Commands correct | Chatter ignored | Agent spoke on chatter | Speech → screen (p50) |
|-------------------|------------------|-----------------|------------------------|-----------------------|
| Clean             | 30 / 31 (96.8%)  | 12 / 12         | 0                      | 3.3 s                 |
| Pink noise -12 dB | 30 / 31 (96.8%)  | 12 / 12         | 0                      | 2.4 s                 |

- Both misses were transcription errors — and the agent did the **safe** thing ("Which table?" / stayed silent).
- Safety cases pass: "Void 7" asks first · unknown item → "Which item?"
- **15 unit tests** ( `npm test` ) · every session recorded and replayable
- ⚠️ Clips are synthesized speech + pink noise (label as such, or record real voices)

## Business

- **Buyer:** independent restaurants and restaurant groups (owner / kitchen manager)
- **Price:** per location / month, optional hardware bundle (headset + tablet)
- **Distribution:** add-on for POS & KDS vendors (Toast, Square, Lightspeed) — the `act()` interface is where their inventory APIs plug in
- **Value:** fewer refunds and comps · faster ticket times · allergy flags that can't be missed · timestamped 86 log from the session timeline
- **Roadmap:** multiple stations · multi-location dashboards · manager analytics

**Thanks! Questions?**

**86 It — say it once. Every screen knows.** <https://86-it.rd5510.workers.dev> · [docs/SUBMISSION.md](#)